

Scaling ablation: ms_per_query vs n_nodes

